

Police as Policymakers: How experiences with policy implementation shape policy representation

Supporting Information

Michael Sierra-Arévalo¹, José J. Alcocer², Lauren Brown³, Raquel Delerme⁴, Brittany Friedman⁴,
Benjamin A.T. Graham⁵, Harry G. Muttram⁶, Jackson Trager⁷, and Nicholas Weller^{*6}

¹Department of Sociology, University of Texas at Austin

²Harvard Law School, Harvard University

³Price School of Public Policy, University of Southern California

⁴Department of Sociology, University of Southern California

⁵Department of Political Science, University of Southern California

⁶Department of Political Science, University of California, Riverside

⁷Department of Psychology, University of Southern California

June 19, 2026

*Corresponding author: nweller@ucr.edu

Contents

A Survey Information A2

A.1 Survey Question Wording A2

A.1.1 Stopped A2

A.1.2 Prior Unfair Treatment A3

A.1.3 Police Legitimacy Index A3

A.1.4 Effectiveness Index A4

A.1.5 Positive Actions Index A5

A.1.6 Legitimate A5

A.1.7 Negative Actions Index A5

A.1.8 Positive Prompts: Topic Modeling A8

A.1.9 Negative Prompts: Topic Modeling A8

A.2 Survey Descriptive Statistics A9

A.3 Additional Index Information A13

B Topic Modeling A15

B.1 Data Preparation and Model Configuration A15

B.2 Topic Consolidation and LLM-Assisted Labeling A16

B.3 Topic Validity A18

B.4 Validity Plots A19

B.5 Significance Testing of Topic Proportions Across Racial Groups A21

C Interview Information A22

C.1 Interview sample and design A22

D Human Subjects Research Ethics A25

The following supporting materials contain descriptive statistics, additional analyses, and other information that are directly referenced in the main text of *Police as Policymakers*.

A Survey Information

The survey was fielded in two waves (November 2023 and June 2024) by Prime Panels, a sampling platform operated by CloudResearch. Prime Panels recruits from approximately 225 partnered market research providers, yielding a global pool of 40–50 million opt-in participants. While not probability-based, the platform’s scale produces samples that approximate U.S. demographics reasonably well ([Chandler et al. 2019](#)).

We provided Prime Panels with a list of Los Angeles city ZIP codes, which the vendor used to target adults 18 and older residing in those areas. Recruitment language varied across supplier channels (email invitations, hub websites, etc.) but was kept generic at the invitation stage (e.g., “Psychology Survey,” “We need your opinions!”) to avoid priming respondents before screening. Invitations included expected length and incentive information.

The combined sample comprises 2,118 respondents: 1,173 from November 2023 and 945 from June 2024. We pool the waves in analyses where appropriate and weight our results back to LA demographics (i.e., age, race, education) using data from the American Community Survey.

A.1 Survey Question Wording

We assess and create various indices described in the main paper. We present here the wording of the questions used to create these indices, as well as the wording of the other primary outcome measures used in our survey.

A.1.1 Stopped

One main variable of interest is whether or not a respondent has been stopped. We used two slightly different questions on both surveys. In the November Wave, we asked:

- Another important goal of our research project is to understand people’s actual experiences with police officers during traffic stops. The following questions ask you for information about your experience(s) during a motor vehicle stop. Have you been stopped by the police

while driving a motor vehicle, NOT including any driving violations captured by camera and ticketed by mail?

- Yes
- No

And in the June Wave, we asked:

- Another important goal of our research project is to understand people’s actual experiences with police officers during traffic stops. The following questions ask you for information about your experience(s) during a motor vehicle stop. Have you been stopped in the last six months by the police while driving a motor vehicle, NOT including any driving violations captured by camera and ticketed by mail?

- Yes
- No

Due to differences in question wording, our main analyses are presented both pooled and broken out by wave. We also treat the June wave as a replication, this time focusing on a more proximate experience.

A.1.2 Prior Unfair Treatment

- Have you ever been treated unfairly by the police?
 - Yes, once
 - Yes, more than once
 - No

This question was only asked in the June Wave.

A.1.3 Police Legitimacy Index

Respondents were asked to indicate their agreement with a series of statements reflecting perceptions of police legitimacy in Los Angeles. Each item was measured on a 5-point Likert scale:

Strongly Disagree, Disagree, Neither Agree nor Disagree, Agree, Strongly Agree. A version of this index was previously validated by [Peyton, Sierra-Arévalo and Rand \(2019\)](#).

- In the next question we'd like to hear more about how you feel about the police in Los Angeles.

Do you agree or disagree with the statements below?

1. They care about the well-being of everyone they deal with
2. They make fair and impartial decisions
3. They act in ways consistent with your own ideas of right and wrong
4. They have the same sense of right and wrong that you do
5. They stand up for values that are important to you
6. They behave according to the law when dealing with people
7. They make decisions based on their biases or opinions
8. They are often dishonest

The last two items are reverse-coded before the creation of the index. These questions were only asked in the June Wave. $\alpha = 0.88$.

A.1.4 Effectiveness Index

Respondents rated police performance on the following dimensions using a 5-point scale: Very poor job, Poor job, Just okay, Good job, Very good job.

- In your opinion, how well are the police doing with:
 1. Dealing with problems that concern your neighborhood
 2. Fighting crime in your neighborhood
 3. Treating people fairly regardless of who they are

$\alpha = 0.89$.

A.1.5 Positive Actions Index

Many items measured respondents' agreement with the following officer behaviors during the traffic stop. Responses were coded on a 5-point Likert scale: Strongly disagree, Disagree, Neither agree nor disagree, Agree, Strongly agree.

- During the traffic stop, the officer:
 1. Comforted and reassured me
 2. Seemed trustworthy
 3. Was fair and evenhanded
 4. Treated me with dignity and respect

$\alpha = 0.93$.

A.1.6 Legitimate

- Would you say the police had a legitimate reason for stopping you?
 - Yes
 - No
 - Don't Know

Note: This question was only asked in the November Wave.

A.1.7 Negative Actions Index

Respondents who reported being stopped by police were asked to indicate which, if any, of a list of negative officer behaviors occurred during the stop (check-all-that-apply, coded 1 if selected and 0 otherwise). Because the June and November instruments used slightly different item sets, we construct three versions of the index: a June-wave version, a November-wave version, and a shared 10-item version that pools items present in both waves. Each index is additive and items are treated as either happening (marked) or not (absent).

November wave (19 items)

- During the stop did the police do any of the following (check all that apply):

1. Shout at you
2. Swear at you
3. Ignore a question you asked
4. Tell you to calm down
5. Interrupt you
6. Threaten to arrest you
7. Threaten to use force against you
8. Point a flashlight at you
9. Push, kick, or grab you
10. Handcuff you
11. Spray you with chemical/pepper spray
12. Use a taser or stungun
13. Use any other type of force
14. Reach into the car
15. Point a gun at you
16. Fire a gun at you

- Additional officer-action items asked separately (coded 1 = Yes, 0 = no):

17. Conducted a frisk
18. Conducted a vehicle search
19. Asked you to exit the car

$\alpha = 0.83$.

June wave (11 items)

- During the stop did the police do any of the following (check all that apply):

1. Shout or swear at you

2. Ignore a question you asked
3. Tell you to calm down
4. Interrupt you
5. Threaten to arrest you
6. Use force against you
7. Handcuff you
8. Use a taser or stungun
9. Reach into the car
10. Point a gun at you
11. Remove you from the car

$\alpha = 0.79$.

For the pooled analysis, we use only items asked in both waves. Since November split a June item into two (e.g., “shout” and “swear,” or “push, kick, or grab” and “use any other type of force”), the shared indicator equals 1 if either component was selected.

- During the stop did the police do any of the following:

1. Shout or swear at you
2. Ignore a question you asked
3. Tell you to calm down
4. Interrupt you
5. Threaten to arrest you
6. Use force against you (push, kick, grab, or other force)
7. Handcuff you
8. Point a gun at you
9. Reach into the car
10. Use a taser or stungun

$\alpha = 0.74$.

A.1.8 Positive Prompts: Topic Modeling

- For the next two questions we would like you to think about an interaction between you and a police officer during a normal traffic stop. For example, imagine you have been pulled over for speeding on a local road. Please answer the following questions about communication with the officer during the interaction.
 - Please describe how you would like the officer to communicate with you. Feel free to include specific words or tone you would like the officer to use.
 - Please describe any actions you would like the officer to take during the traffic stop.

A.1.9 Negative Prompts: Topic Modeling

- As a reminder, in answering the next two questions we would like you to think about an interaction between you and a police officer during a normal traffic stop. For example, imagine you have been pulled over for speeding on a local road.
 - Please describe how you would NOT like the officer to communicate with you. Feel free to discuss specific words or tone you would not want the officer to use.
 - Please describe any actions that you would NOT want the officer to take during the traffic stop.

A.2 Survey Descriptive Statistics

Table A2 presents the basic descriptive statistics for the combined survey waves. Table A3 presents the basic descriptive statistics for the November Wave. Table A4 presents the basic descriptive statistics for the June Wave.

Table 1: Weighted Descriptive Statistics: Both Waves

Variable	N = 2,118 ¹
Race	
White	583 (28%)
Black	189 (8.9%)
Hispanic	973 (46%)
Asian	280 (13%)
Other	93 (4.4%)
Income	
Less than \$20,000	431 (20%)
\$20,000 to \$40,000	460 (22%)
\$40,000 to \$60,000	387 (18%)
\$60,000 to \$80,000	288 (14%)
\$80,000 to \$100,000	200 (9.4%)
\$100,000 to \$150,000	212 (10%)
More than \$150,000	139 (6.6%)
Education	
High school or less	862 (41%)
Some college	532 (25%)
Bachelor's degree or higher	724 (34%)
Age	
18-24 years old	254 (12%)
25-34 years old	468 (22%)
35-44 years old	394 (19%)
45-54 years old	326 (15%)
55-64 years old	303 (14%)
65+ years old	373 (18%)
Stopped	822 (39%)
Positive Actions Index	12.1 (4.7)
Unknown	1,296
Negative Actions Index (Shared)	1.1 (1.7)
Effectiveness Index	9.0 (3.1)
Stop Legitimate?	425 (67%)

¹n (%); Mean (SD)

Table 2: Weighted Descriptive Statistics November Wave

Variable	N = 1,173 ¹
Race	
White	323 (28%)
Black	105 (8.9%)
Hispanic	539 (46%)
Asian	155 (13%)
Other	52 (4.4%)
Income	
Less than \$20,000	240 (20%)
\$20,000 to \$40,000	237 (20%)
\$40,000 to \$60,000	194 (17%)
\$60,000 to \$80,000	178 (15%)
\$80,000 to \$100,000	121 (10%)
\$100,000 to \$150,000	114 (9.7%)
More than \$150,000	90 (7.7%)
Education	
High school or less	477 (41%)
Some college	294 (25%)
Bachelor's degree or higher	401 (34%)
Age	
18-24 years old	141 (12%)
25-34 years old	259 (22%)
35-44 years old	218 (19%)
45-54 years old	181 (15%)
55-64 years old	168 (14%)
65+ years old	206 (18%)
Stopped	679 (58%)
Positive Actions Index	12.2 (4.6)
Negative Actions Index (November Items)	1.8 (2.7)
Effectiveness Index	9.0 (3.0)
Stop Legitimate?	424 (67%)

¹n (%); Mean (SD)

Table 3: Weighted Descriptive Statistics June Wave

Variable	N = 945 ¹
Race	
White	260 (28%)
Black	84 (8.9%)
Hispanic	434 (46%)
Asian	125 (13%)
Other	42 (4.4%)
Income	
Less than \$20,000	194 (20%)
\$20,000 to \$40,000	220 (23%)
\$40,000 to \$60,000	195 (21%)
\$60,000 to \$80,000	112 (12%)
\$80,000 to \$100,000	79 (8.4%)
\$100,000 to \$150,000	97 (10%)
More than \$150,000	49 (5.2%)
Education	
High school or less	385 (41%)
Some college	237 (25%)
Bachelor's degree or higher	323 (34%)
Age	
18-24 years old	113 (12%)
25-34 years old	209 (22%)
35-44 years old	176 (19%)
45-54 years old	146 (15%)
55-64 years old	135 (14%)
65+ years old	166 (18%)
Stopped	145 (15%)
Positive Actions Index	11.5 (5.1)
Negative Actions Index (June Items)	2.2 (2.4)
Effectiveness Index	9.0 (3.1)
Legitimacy Index	23.7 (6.2)
Treated Unfairly by Police?	
No	553 (59%)
Yes, once	199 (21%)
Yes, more than once	192 (20%)

¹n (%); Mean (SD)

A.3 Additional Index Information

We use several indices in this paper. Figure 1 presents the alpha (survey-weighted), percentage of total item variance explained by the first principal component, and the resulting scree plot for each (i.e., the eigenvalues for each component). Our main indices, the Positive Actions Index, the Effectiveness Index, and the Legitimacy Index, are all plausibly unidimensional. The Negative Actions Indices have high alphas but do not meet conventional metrics for unidimensionality. That said, (1) the first PC has an eigenvalue magnitudes higher than the second PC, and (2) we do not use these indices as main results; rather, they supplement our main findings. Because we use the negative action measures descriptively to supplement our main findings, we retain the additive form as a crude summary rather than constructing separate subscales




Figure 1: Scree Plots by Index

Main Figures Pooling Black and Hispanic

In our manuscript, we compare the experiences and attitudes of Black and Hispanic individuals to those of Whites. To ensure our results are not driven by a lack of power, we pool these nonwhite groups together and present all the main plots again in Figure 2. All findings hold when using this

alternative operation.




Figure 2: All Main Figures are Robust to Pooling Black and Hispanic Respondents.

B Topic Modeling

B.1 Data Preparation and Model Configuration

Open-ended survey responses were analyzed separately for the positive and negative question conditions. Prior to topic modeling, verbal and action-based responses within each condition were concatenated into a single text field per respondent, yielding two primary corpora: one capturing preferences for desired officer behavior and one capturing preferences for undesired officer behavior.

Topic modeling was implemented using [BERTopic](#), facilitated through [TextViz Studio](#), a Python-based text analysis application developed by one of the coauthors that automates complex natural language processing pipelines while allowing modular configurations. The BERTopic pipeline proceeded through four sequential stages.

Document Embedding. Each response was converted into a dense vector representation using the [bge-small-en-v1.5](#) sentence transformer model, a pre-trained English-language embedding model that encodes the full contextual meaning of each response as a fixed-length vector in a high-dimensional semantic space. Responses were embedded in their raw, unprocessed form, consistent with best practices for transformer-based models that require complete syntactic context to generate accurate semantic representations.

Dimensionality Reduction. Document embeddings were reduced from their native dimensionality to five dimensions using Uniform Manifold Approximation and Projection (UMAP; [McInnes, Healy and Melville 2018](#)), which preserves both local and global structure in the high-dimensional data while producing a lower-dimensional representation suitable for density-based clustering. A fixed random seed of 44 was applied to ensure full reproducibility of all downstream outputs.

Clustering. Dimensionality-reduced embeddings were clustered using Hierarchical Density-Based Spatial Clustering of Applications with Noise (HDBSCAN; [Campello et al. 2015](#); [McInnes, Healy and Astels 2017](#)), the default clustering algorithm in BERTopic. HDBSCAN identifies clusters by locating regions of higher density in the embedding space, treating densely packed groups of semantically similar responses as topics and sparser surrounding points as noise designated Topic -1. Each response is assigned to the single dense region its embedding falls within. The number

of topics was determined automatically by HDBSCAN without any predetermined target, allowing the density structure of the data to govern the granularity of the solution.

Following initial clustering, outlier reduction was applied at a class-based TF-IDF similarity threshold of 0.1 , followed by a probability distributions pass, to reassign responses that HDBSCAN had designated as noise to their nearest substantive topic.

Topic Representation. For each identified cluster, an initial set of candidate keywords was generated using the class-based Term Frequency-Inverse Document Frequency procedure (c-TF-IDF; [Grootendorst 2022](#)) with English stop words removed. A CountVectorizer with an n-gram range of one to three was applied to capture both individual terms and short phrases. Candidate keywords were then re-ranked using the KeyBERTInspired representation model, which scores each candidate term by computing the cosine similarity between its embedding and the embeddings of the documents within the topic, retaining those terms that are most semantically representative of the topic’s content rather than merely the most frequent ([Grootendorst 2022](#)). This two-stage approach combines the discriminative power of c-TF-IDF, which identifies terms that are distinctive to a topic relative to all other topics, with the semantic grounding of embedding-based keyword selection, producing topic keywords that are both topically distinctive and semantically coherent.

The initial topic solution yielded 15 topics for the positive condition and 25 topics for the negative condition.

B.2 Topic Consolidation and LLM-Assisted Labeling

Following initial topic extraction, the resulting topics were reviewed qualitatively and consolidated through a principled merging procedure. The review was guided by three sources of evidence weighted in the following order of priority. Primary weight was given to close reading of each topic’s keywords and representative documents, which together constitute the most direct evidence of a topic’s substantive content. Secondary weight was given to the intertopic distance map, which was used as a heuristic to identify vocabulary-based proximity between topics. Tertiary and minimal weight was given to the hierarchical clustering dendrogram produced by BERTopic, which was treated as structural context rather than a primary decision criterion, given that all topics were expected to connect at some level of the hierarchy due to the primed and domain-constrained nature

of the survey responses.

After topics were merged manually, topic labels and descriptions were generated using GPT-5.5. For each merged topic, the model was provided with the pooled keywords and representative documents from all constituent topics, and prompted to produce a short label of five words or fewer, and a two-to-three sentence description of the core theme. The prompt explicitly instructed the model to ground its label in the substantive content of the keywords and representative responses rather than surface vocabulary patterns shared across topics. To minimize label redundancy across topics, A sequential labeling procedure was adopted in which each successive prompt included the labels already assigned to prior topics, with an explicit instruction that the new label must be meaningfully distinct and must not reproduce a synonym or minor term rewording of any previously assigned label. The full prompt structure used for labeling is reproduced below:

```
You are analyzing open-ended survey responses where participants described
the verbal and action-based behaviors they [would WANT / would NOT WANT] to
experience during a police traffic stop.
```

```
The topic is described by these keywords: {keywords}.
```

```
The following are representative survey responses belonging to this topic:
{representative_docs}.
```

```
Your label must reflect which of these the topic primarily represents, not
just the surface vocabulary shared across topics. Based solely on the keywords
and representative responses above, provide: (1) a short, precise topic label
of five words or fewer, and (2) a two-to-three sentence description of the
core theme. Respond in exactly this format with no additional text.
```

```
LABEL: [label here]
```

```
DESCRIPTION: [description here].
```

For topics after the first, the following instruction was appended:

“The following labels have already been assigned to other topics in this dataset. Your label must be meaningfully distinct from all of them—do not produce a synonym or minor rewording of any label below: [existing labels].”

B.3 Topic Validity

Vocabulary-Based Distinctiveness. For each merged topic, we constructed a combined vocabulary profile by summing the c-TF-IDF vectors of all constituent initial topics and L2-normalizing the result. These normalized vectors were projected to two dimensions using seeded UMAP to produce an intertopic distance map in which each topic is represented as a bubble whose position reflects its vocabulary profile relative to all other topics and whose size is proportional to its document count. Topics that are spatially separated in this projection use characteristically different vocabulary to represent their content. Figures 3 and 4 present these maps for the positive and negative conditions, respectively. In both figures, topics occupy visually distinct regions of the projected space. Where visual overlap occurs between topic bubbles, this is an artifact of bubble sizing from document counts rather than topical similarity—the position of each bubble’s center represents the topic’s location in the vocabulary space.

Semantic Distinctiveness. To assess distinctiveness in the semantic embedding space independently of vocabulary, we computed a sentence embedding centroid for each topic by averaging the document embeddings of all responses assigned to that topic using the same `bge-small-en-v1.5` model applied during BERTopic fitting. These centroids were L2-normalized and projected to two dimensions via seeded UMAP to produce a parallel intertopic distance map in which each topic’s position reflects the average semantic meaning of its constituent responses. Figures 5 and 6 present these maps for the positive and negative conditions, respectively. Topics again occupied visually distinct regions of the projected space, providing convergent evidence from an independent representation that the topic solution captures substantively different themes.

Taken together, the spatial separation observed in both projections indicates that the consolidated topic structures are not arbitrary partitions of a homogeneous response corpus but reflect genuine thematic differentiation in what respondents expressed. It should be noted that the absolute positions and distances between topics in UMAP projections are not directly interpretable, as UMAP optimizes for preserving local neighborhood structure rather than precise global distances. Thus, the figures should be read as evidence of relative separation rather than as precise distance measurements.

B.4 Validity Plots




Figure 3: Vocabulary Profile of Positive Topics




Figure 4: Vocabulary Profile of Negative Topics




Figure 5: Semantic Profile of Positive Topics




Figure 6: Semantic Profile of Negative Topics

B.5 Significance Testing of Topic Proportions Across Racial Groups

To assess whether the weighted proportion of respondents assigned to each topic differed significantly across racial groups, we estimated a series of weighted logistic regressions with topic membership as the binary outcome and racial group as the predictor. For each topic k , a binary indicator was constructed taking the value of 1 if a respondent was assigned to topic k and 0 otherwise. A separate model was estimated for each topic:

$$\log\left(\frac{p_{ik}}{1-p_{ik}}\right) = \beta_0 + \beta_1\text{Black}_i + \beta_2\text{Hispanic}_i + \beta_3\text{Asian}_i + \beta_4\text{Other}_i + \epsilon_i \quad (1)$$

where p_{ik} is the probability that respondent i is assigned to topic k and White respondents serve as the reference category. Exponentiated coefficients (odds ratios) and their 95% confidence intervals are reported in Tables 4 and 5 for the positive and negative conditions, respectively.

Table 4: Odds Ratios from Weighted Logistic Regressions Predicting Positive Topic Membership by Race (Reference Group: White)

Topic	Asian	Black	Hispanic
Calm respectful communication	0.77 [0.52, 1.16]	0.89 [0.56, 1.45]	1.09 [0.80, 1.48]
Explain stop reason clearly	0.79 [0.44, 1.36]	0.80 [0.40, 1.48]	0.75 [0.51, 1.12]
Friendly greeting and wave	1.30 [0.25, 5.63]	1.15 [0.14, 5.91]	1.11 [0.37, 3.82]
Keep weapon holstered	0.29 [0.01, 1.78]	0.83 [0.10, 3.77]	0.67 [0.23, 2.03]
No specific preference	3.74*** [1.73, 8.48]	2.76** [1.08, 6.92]	1.67 [0.83, 3.60]
Polite professional tone	1.65 [0.54, 4.81]	1.02 [0.20, 3.78]	1.07 [0.45, 2.72]

* $p < 0.10$ ** $p < 0.05$ *** $p < 0.01$

Values presented as odds ratio [95% confidence interval].

The Other racial group category is excluded from this table due to small cell counts.

Table 5: Odds Ratios from Weighted Logistic Regressions Predicting Negative Topic Membership by Race (Reference Group: White)

Topic	Asian	Black	Hispanic
Abusive yelling or insults	0.95 [0.62, 1.45]	0.58* [0.31, 1.00]	0.85 [0.63, 1.17]
Aggression or disrespect	1.21 [0.91, 1.61]	0.98 [0.70, 1.36]	0.95 [0.77, 1.17]
Forceful shouting or demands	0.93 [0.54, 1.54]	1.61* [0.95, 2.66]	0.81 [0.56, 1.19]
Hostile threatening conduct	0.73 [0.37, 1.35]	1.45 [0.78, 2.59]	1.22 [0.82, 1.85]
No specific behavior identified	1.15 [0.63, 2.05]	1.24 [0.62, 2.34]	1.13 [0.74, 1.75]
Profiling or baseless arrest	0.72 [0.17, 2.31]	0.12 [0.00, 1.12]	1.79 [0.89, 3.95]
Rude ticketing	1.14 [0.46, 2.63]	0.86 [0.25, 2.36]	0.62 [0.30, 1.28]
Weapons or forced removal	0.83 [0.56, 1.21]	0.93 [0.60, 1.43]	1.15 [0.89, 1.50]

* $p < 0.10$ ** $p < 0.05$ *** $p < 0.01$

Values presented as odds ratio [95% confidence interval].

The Other racial group category is excluded from this table due to small cell counts.

C Interview Information

C.1 Interview sample and design

Our sample of 30 semi-structured interviews was collected between July 2023 and September 2024. Interviews were conducted in person or via online video conference and were recorded with the consent of participants, then transcribed and de-identified with a unique code and pseudonym. This study’s qualitative analysis draws on semi-structured interviews with 30 people who live or work in the city of Los Angeles. Interviewees’ pseudonyms, self-reported race, gender, and age are displayed in Table 6.

Our sample of interviewees is non-random and differs from the demographic composition of Los Angeles in some ways. Relative to city demographics ([U.S. Census Bureau 2022](#)), our interview sample has an underrepresentation of Latinos and Asians, and an overrepresentation of Black respondents. The proportion of White respondents in our sample accords with city-level demographics. Of course, the analytic value of interviews is not predicated on their generalizability to a population based on the logic of random sampling. Instead, the value of qualitative methods and data often lies in their ability to elucidate mechanisms—the “how” and “why”—which underlie social phenomena ([Small 2009](#)).

The majority of our interview sample is, by design, composed of non-White people who are more likely to have vicarious and direct experiences with police informed by the intensive policing of many minority communities (Prowse, Weaver and Meares 2020). This sample does not support (and we make no efforts to make) confident claims of the distribution of interviewees’ experiences or statistical assessment of variation across interviewees. Instead, we compare self-reported experiences with and expectations of police across interviewees to describe if, how, and why interviewees’ responses vary along race, gender, and other salient characteristics.

To capture individuals’ experiences interacting with police, we constructed our interview guide to focus on the content of respondents’ prior contact with police and their view of police. Building on prior work that estimates how the content of a single police interaction affects individuals’ perceptions of police (Dobson, Dittmann and Yeager 2025; Peyton, Sierra-Arévalo and Rand 2019), our interviews are designed flexibly to cover prior contacts that inform how people relate to, interact with, or avoid police (Brunson 2007; Deckard et al. 2025; Rosenbaum et al. 2005). Unlike instruments which rely on predefined response choices, open-ended interview questions allow respondents to express in their own words how they experienced contact with police, as well the degree to which prior contact informs their contemporary view of police.

Our questions about respondents’ prior experiences with police include specific questions about their experiences as the driver or passenger during traffic stops. We explore individuals’ expectations—what they hoped or desired might happen during a stop, as well as what they worried or feared might occur—during traffic stop interactions. We also asked respondents about officers’ behavior during stops and how respondents were affected by those behaviors. We focus in particular on two actions that LAPD officers are explicitly trained to take early in traffic stop interactions: 1) introducing themselves to vehicle occupants, and 2) explaining the reason for the stop. By asking respondents about common actions officers may take early in a stop, we can better understand how respondents understand, experience, and respond to these behaviors. In conjunction with questions about respondents’ attitudes about police and their prior experiences with police, our interviews enable a more granular view of how views of police and the experience of police interactions are, rather than temporally discrete phenomena, informed by the content and experience of past interactions with police.

Table 6: Demographic Information of Interviewees

Name	Race/Ethnicity	Gender	Age
Amelia	Asian	Female	31
Jennifer	Asian	Female	35
Oliver	Asian	Male	Unreported
Emma	Asian	Female	Unreported
Olivia	Black	Female	32
Sophia	Black	Female	59
Samuel	Black	Male	28
Elijah	Black	Male	34
Isabella	Latino	Female	33
Victoria	Latino	Female	53
Scarlett	Latino	Female	23
Henry	Latino	Male	37
Chloe	Other/Mixed Race	Female	35
Leo	Other/Mixed Race	Male	Unreported
William	Black	Male	19
Audrey	Black	Female	20s
Marc	Black	Male	Unreported
Grace	White	Female	29
Kelsie	White	Female	23
Evelyn	White	Female	34
Ella	White	Female	44
Ava	White	Female	54
Emily	White	Female	27
Madison	White	Female	44
Anna	White	Female	29
Sarah	White	Female	25
Liam	White	Male	33
Noah	White	Male	38
James	White	Male	36
Corey	Black	Male	31

D Human Subjects Research Ethics

This paper relies on survey responses from over 2000 individuals and in-depth interviews from a smaller sample. Approval was secured in advance from the University of Southern California Internal Review Board. Due to the minimal risks posed to respondents, both the survey and the interview research components of the study were classified as “exempt.”

To preserve subject anonymity, no names were collected in the survey or recorded in interview notes. In instances where interviews were recorded, transcripts were de-identified prior to analysis. All names used for interviewees in the write-up are pseudonyms.

References

- Brunson, Rod K. 2007. ““Police don’t like black people”: African-American young men’s accumulated police experiences.” *Criminology & public policy* 6(1):71–101.
- Campello, Ricardo J. G. B., Davoud Moulavi, Arthur Zimek and Jörg Sander. 2015. “Hierarchical Density Estimates for Data Clustering, Visualization, and Outlier Detection.” *ACM Trans. Knowl. Discov. Data* 10(1).
URL: <https://doi.org/10.1145/2733381>
- Chandler, Jesse, Cheskie Rosenzweig, Aaron J Moss, Jonathan Robinson and Leib Litman. 2019. “Online panels in social science research: Expanding sampling methods beyond Mechanical Turk.” *Behavior research methods* 51(5):2022–2038.
- Deckard, Faith M, Shannon Malone Gonzalez, Yasmiyn Irizarry and Jaime Feng-Yuan Hsu. 2025. “Perpetual encounters: reconceptualizing police contact and measuring its relationship to black women’s mental health.” *Social Forces* p. soaf070.
- Dobson, Kyle SH, Andrea G Dittmann and David S Yeager. 2025. “A transparency statement improves trust in community-police interactions.” *Nature communications* 16(1):2285.
- Grootendorst, Maarten. 2022. “BERTopic: Neural topic modeling with a class-based TF-IDF procedure.” *arXiv preprint arXiv:2203.05794* .
- McInnes, Leland, John Healy and James Melville. 2018. “Umap: Uniform manifold approximation and projection for dimension reduction.” *arXiv preprint arXiv:1802.03426* .
- McInnes, Leland, John Healy and Steve Astels. 2017. “hdbscan: Hierarchical density based clustering.” *Journal of Open Source Software* 2(11):205.
URL: <https://doi.org/10.21105/joss.00205>
- Peyton, Kyle, Michael Sierra-Arévalo and David G Rand. 2019. “A field experiment on community policing and police legitimacy.” *Proceedings of the national Academy of sciences* 116(40):19894–19898.

- Prowse, Gwen, Vesla M Weaver and Tracey L Meares. 2020. "The state from below: Distorted responsiveness in policed communities." *Urban Affairs Review* 56(5):1423–1471.
- Rosenbaum, Dennis P, Amie M Schuck, Sandra K Costello, Darnell F Hawkins and Marianne K Ring. 2005. "Attitudes toward the police: The effects of direct and vicarious experience." *Police quarterly* 8(3):343–365.
- Small, Mario Luis. 2009. "How many cases do I need?" On science and the logic of case selection in field-based research." *Ethnography* 10(1):5–38.
- U.S. Census Bureau. 2022. "American Community Survey 2022: QuickFacts Table for Los Angeles City, California." <https://www.census.gov/quickfacts/fact/table/losangelescacitycalifornia/PST045224>. Data derived from the American Community Survey (U.S. Census Bureau).